

What are we testing?

Can a coding agent build a working forecast from scratch, and how does that depend on the model?

Does adding Rhiza weather skills or ACCORD libraries improve completion, time or cost?

When the next forecast is needed, does retained code and fitted state reduce the work?

Report completion, forecast error, tokens, time and cost separately. Treat the expected advantages as hypotheses.

The three arms we actually tested

Scratch

Standard scientific Python; the agent builds the workflow.

Rhiza skills

The same Python environment + Rhiza weather skills catalog and scripts.

ACCORD

The same Python environment + africas2s (DeepScale) and Rosetta/acmadDL, with API documentation.

Identical supplied data and budgets. Every arm can retain its own code, inputs and fitted models.

Observed use: Rhiza catalog context, no script calls; ACCORD agents called africas2s. Data retrieval was not tested.