

Can a coding agent build a seasonal forecast workflow?

September-issued OND rainfall in a Kenyan study box.

Aggregate monthly inputs, fit a calibration, save the fitted model, and produce a new forecast.

Scratch model	Checks passed	Tokens	Cost	Time
Fable 5.1	0/12	236,218	\$4.38	9.6 min
Haiku 4.5	6/12	473,419	\$0.60	4.3 min

Neither scratch model completed its initial attempt. Fable later reached 11/12 checks using saved work; this pilot does not establish a model capability threshold.

Same inputs and attempt limits. One attempt per condition: an illustration, not a general model ranking.

Monthly data was supplied; live data discovery and retrieval were not tested.

Do skills and domain libraries reduce the work?

Same Fable 5.1 model, task, normalized inputs, resources and limits. Only the available domain tools differ.

Arm	Checks	Tokens	Cost	Time	Test RPS
Python only	0/12	236,218	\$4.38	9.6 min	—
Python + Rhiza	11/12	327,196	\$4.28	5.7 min	0.278
Python + ACCORD	0/12	341,771	\$4.51	6.0 min	—

Rhiza reached 11/12 checks on the first attempt. ACCORD inspected the libraries but wrote its own calibration and did not complete. No library advantage established.

RPS is lower-is-better; two previously inspected development years, not untouched scientific validation.

Numerical checks, offline replay and an evidence-citing LLM judge assess the outputs.

What survives into the next forecast?

Fresh conversation, same arm’s saved code, fitted model, data and handoff notes. Apply to a different forecast batch.

Arm	Saved state	Follow-up checks	Initial → follow-up cost	Follow-up time
Python only	code only	11/12	\$4.38 → \$3.65	5.7 min
Python + Rhiza	model.nc	11/12	\$4.28 → \$4.20	3.4 min
Python + ACCORD	code only	1/12	\$4.51 → \$4.00+	10.0 min

Rhiza’s frozen predictor ran the next batch in 1.5 s with zero LLM tokens. Conversational cost barely fell (\$4.28 → \$4.20); scratch recovered, ACCORD stayed incomplete.

All saved code, models, notes, figures and usage are retained. + marks unrecorded usage from a timed-out request.

No fresh-workspace controls; recovery from an incomplete initial attempt is not reuse of a working forecast.